


RF-LISSOM CUDA

- Giacomo Spigler

- An RF-LISSOM simulator which takes advantage of CUDA compatible GPUs.
- Up to 5x (10x peak) performance gain on GT200 GPUs
- Tesla compatible system (as many GPUs as you want can be used to speed up the simulation even more)

RF-LISSOM

- RF-LISSOM : Laterally Interconnected Synergetically Self-Organizing Map
- This special kind of SOM introduces lateral connections (short-range excitatory and long-range inhibitory) into the classical Self-Organizing Map framework, producing an accurate simulation of the neural cortex at neural columns level.
- This model proved successful in simulating visual phenomena occurring in Primary Visual Cortex V1, among whom we can list development of Orientation Selectivity, Ocular Dominance, Depth Perception, and a direct observation of the Tilt-After Effect.



-Moreover, the model can probably simulate visual processing in the brain up to the Object Recognition level, to Visual Segmentation and to Perceptual Grouping.

-The image on the left shows the typical connections being simulated by the model.

RF-LISSOM

-RF-LISSOM basic operations are quite easy to describe, but they actually need an enormous amount of memory and of computational power.

-RF-LISSOM can be divided into 3 steps: First Activation, Settling and Weights Adjusting.

-First Activation : Once an input is presented on the input layer, every LISSOM's neuron computes a first activation based on its current Afferent Weights. The result is then processed with a piecewise linear approximation of the sigmoidal function.

-Settling : This phase consists of running a few (typical 9-13) steps, which will simulate lateral connections, and which will produce some stable blobs of neural activity, representing a stable state of the network (given the current input).

$$\eta_{ij}(t) = \sigma \left(s_{ij} + \gamma_E \sum_{kl} \eta_{kl}(t-1) E_{kl,ij} - \gamma_I \sum_{kl} \eta_{kl}(t-1) I_{kl,ij} \right)$$

$\eta(t)$ is the firing strength of neuron (i, j), and the 3 activations inside the sigmoidal function are: First Activation for neuron (i, j), Excitatory Activation with neuron (k, l) and Inhibitory Activation (neuron (k, l)).

RF-LISSOM

-Weights Adjusting : Weights are then modified according to the last neural activation recorded using Hebbian Rule.

They are then normalized to sum to 1; hence, the modification of a single weight it's compared to those of the other weights (of the same type, eg, afferent).

Every weight, on weights adjusting phase, either increases or decreases its **relative** strenght, gathering more importance.

$$w'_{pq,ij} = \frac{w_{pq,ij} + \alpha X_{pq} \eta_{ij}}{\sum_{uv} (w_{uv,ij} + \alpha X_{uv} \eta_{ij})}$$

CUDA Simulation

-Here are shown the main features of the CUDA implementation:

- Every layer contains three objects, called Projections, representing its different weights types: Afferent, Excitatory and Inhibitory

- Every Projection's function is almost fully interexchangeable with the other ones (eg, adjusting afferent weights is accomplished with the same code as used to adjust inhibitory weights; the only things we change are parameters).


- RF-LISSOM is intrinsically a parallel problem, as Neural Networks always are; hence, the highly parallel architecture of CUDA and CUDA compatible GPUs lead to a great improvement in speed.

- The main bottleneck is memory: every iteration consists of tens of thousands, and perhaps millions of sparse access to weights in memory, which prevents from coalescing memory reads and writes.


- A partial solution to this problem has been found using Textures bounded to linear memory for fast, cached reads, and memory writes were reduced.

Examples

- After training on synthetic data (elongated gaussians), we can test both what the network looks like and how it behaves.
- Here is the network's activity when an input is presented (shown at the left) before training.




- And this is what it looks like after Self-Organization occurred (please notice the formation of little blobs, which are tuned to that specific orientation and won't respond to any other stimulus).




Examples

-Now, let's have a look at LISSOM's weights after 20000 training steps.

-Excitatory Connections are shrunk during training:



-Inhibitory Connections develop among the neuron's preferred orientation, and help producing the final stable blobs.



-Afferent Connections clearly show their preferred orientation.



Conclusions and Future Work

- Research on RF-LISSOM and on its CUDA implementation will lead to important results in understanding how the visual cortex, and probably the whole neural cortex work.
- Next improvements will be:
 - Stable support for multiGPU system (it's currently working, but it needs some fixing)
 - Support for CUDA Clusters (clusters whose nodes are embedded with CUDA compatible hardware), through MPI.
 - Testing and Research on many LISSOM layers connected into a hierarchy.
 - Object Recognition and research on how the neural cortex processes complex data from the real world, producing stable representations of the environment.
 - We are currently investigating some tests to study Tommaso Poggio's Standard Model with a Model of the Visual Cortex produced by a synthesis of RF-LISSOM and current knowledge of neuroanatomy of this portion of the brain.
- In conclusion, CUDA allows us to process increasingly bigger areas of the brain in much less time than what it would take on a classical CPU, and costs are kept down: I am actually an high school student, and I have no source of fundings. However, with CUDA, I can gather access to the computational power I need!

References

- 1) ``Computational Maps in the Visual Cortex``, Miikkulainen, Bednar, Choe, and Sirosh (New York: Springer 2005) <http://nn.cs.utexas.edu/computationalmaps/>
- 2) RF-LISSOM: http://homepages.inf.ed.ac.uk/jbednar/rflissom_small.html
- 3) CUDACluster: <http://cudacluster.nvidia.com>
- 4) nVidia: <http://www.nvidia.com>
- 5) CUDA: http://www.nvidia.com/object/cuda_home.html
- 6) Poggio's CBCL: <http://cbcl.mit.edu>
- 7) Poggio's Standard Model: <http://cbcl.mit.edu/projects/cbcl/publications/ai-publications/2005/AIM-2005-036.pdf>

RF-LISSOM CUDA

END